



Installing Black Duck using OpenShift

Black Duck 2022.10.3

Contents

- Preface.....3**
 - Black Duck documentation..... 3
 - Customer support..... 3
 - Synopsys Software Integrity Community.....4
 - Training.....4
 - Synopsys Statement on Inclusivity and Diversity.....4
- 1. Installing using Synopsysctl..... 5**
- 2. Hardware requirements.....6**
- 3. PostgreSQL versions..... 9**
 - General Migration Process..... 9
- 4. Migrating on OpenShift with Helm.....10**

Preface

Black Duck documentation

The documentation for Black Duck consists of online help and these documents:

Title	File	Description
Release Notes	release_notes.pdf	Contains information about the new and improved features, resolved issues, and known issues in the current and previous releases.
Installing Black Duck using Docker Swarm	install_swarm.pdf	Contains information about installing and upgrading Black Duck using Docker Swarm.
Getting Started	getting_started.pdf	Provides first-time users with information on using Black Duck.
Scanning Best Practices	scanning_best_practices.pdf	Provides best practices for scanning.
Getting Started with the SDK	getting_started_sdk.pdf	Contains overview information and a sample use case.
Report Database	report_db.pdf	Contains information on using the report database.
User Guide	user_guide.pdf	Contains information on using Black Duck's UI.

The installation methods for installing Black Duck software in a Kubernetes or OpenShift environment are Synopsysctl and Helm. Click the following links to view the documentation.

- [Helm](#) is a package manager for Kubernetes that you can use to install Black Duck.
- [Synopsysctl](#) is a cloud-native administration command-line tool for deploying Black Duck software in Kubernetes and Red Hat [OpenShift](#).

Black Duck integration documentation is available on [Confluence](#).

Customer support

If you have any problems with the software or the documentation, please contact Synopsys Customer Support.

You can contact Synopsys Support in several ways:

- Online: <https://www.synopsys.com/software-integrity/support.html>
- Phone: See the Contact Us section at the bottom of our [support page](#) to find your local phone number.

To open a support case, please log in to the Synopsys Software Integrity Community site at <https://community.synopsys.com/s/contactsupport>.

Another convenient resource available at all times is the [online customer portal](#).

Synopsys Software Integrity Community

The Synopsys Software Integrity Community is our primary online resource for customer support, solutions, and information. The Community allows users to quickly and easily open support cases and monitor progress, learn important product information, search a knowledgebase, and gain insights from other Software Integrity Group (SIG) customers. The many features included in the Community center around the following collaborative actions:

- **Connect** – Open support cases and monitor their progress, as well as, monitor issues that require Engineering or Product Management assistance
- **Learn** – Insights and best practices from other SIG product users to allow you to learn valuable lessons from a diverse group of industry leading companies. In addition, the Customer Hub puts all the latest product news and updates from Synopsys at your fingertips, helping you to better utilize our products and services to maximize the value of open source within your organization.
- **Solve** – Quickly and easily get the answers you're seeking with the access to rich content and product knowledge from SIG experts and our Knowledgebase.
- **Share** – Collaborate and connect with Software Integrity Group staff and other customers to crowdsource solutions and share your thoughts on product direction.

[Access the Customer Success Community](#). If you do not have an account or have trouble accessing the system, click [here](#) to get started, or send an email to community.manager@synopsys.com.

Training

Synopsys Software Integrity, Customer Education (SIG Edu) is a one-stop resource for all your Black Duck education needs. It provides you with 24x7 access to online training courses and how-to videos.

New videos and courses are added monthly.

At Synopsys Software Integrity, Customer Education (SIG Edu), you can:

- Learn at your own pace.
- Review courses as often as you wish.
- Take assessments to test your skills.
- Print certificates of completion to showcase your accomplishments.

Learn more at <https://community.synopsys.com/s/education> or for help with Black Duck, select **Black Duck**

Tutorials from the Help menu () in the Black Duck UI.

Synopsys Statement on Inclusivity and Diversity

Synopsys is committed to creating an inclusive environment where every employee, customer, and partner feels welcomed. We are reviewing and removing exclusionary language from our products and supporting customer-facing collateral. Our effort also includes internal initiatives to remove biased language from our engineering and working environment, including terms that are embedded in our software and IPs. At the same time, we are working to ensure that our web content and software applications are usable to people of varying abilities. You may still find examples of non-inclusive language in our software or documentation as our IPs implement industry-standard specifications that are currently under review to remove exclusionary language.


1. Installing using Synopsysctl


OpenShift™ is an orchestration tool from Red Hat used for managing cloud workloads through containers.

Synopsysctl is a command line tool that assists in the deployment and management of Synopsys software in Kubernetes and OpenShift clusters. After synopsysctl is installed, you can leverage it to easily deploy and manage Synopsys software.

Click [here](#) for documentation about installing and using synopsysctl.

If you are a customer using Kubernetes and are using an install method other than synopsysctl, please contact Synopsys Customer Support for migration assistance.

 **Note:** For scalability sizing guidelines, see the Container Scalability section of the Black Duck Release Notes.

 **CAUTION:** Do not delete data from the Black Duck database (bds_hub) unless directed to do so by a Synopsys Technical Support representative. Be sure to follow appropriate backup procedures. Deletion of data will cause errors ranging from UI problems to complete failure of Black Duck to start. Synopsys Technical Support cannot recreate deleted data. Synopsys will provide support on a best-effort basis if no backups are available.


2. Hardware requirements

The performance data below was gathered using Black Duck 2022.10.0 with reduced signature scan persistence (default) and Synopsys Detect 8.0.0. SPH values are calculated using a mix of signature scans, package manager detector scans and rapid scans. Average scan sizes vary from customer to customer so exact SPH throughput is highly customer specific. These metrics were gathered from Google Cloud Platform, which provides different database read/write IOPS for different configurations.

10sph	Scans per hour: 50 SPH % Increase: 400% APIs per hour: 2,500 Project Versions: 10,000	IOPS: <ul style="list-style-type: none"> Read: 15,000 Write: 9,000 Black Duck Services: <ul style="list-style-type: none"> CPU: 12 core Memory: 30 GB PostgreSQL: <ul style="list-style-type: none"> CPU: 2 core Memory: 8 GB 	Total: <ul style="list-style-type: none"> CPU: 14 core Memory: 38 GB
120sph	Scans per hour: 120 SPH % Increase: 0% APIs per hour: 3,000 Project Versions: 13,000	IOPS: <ul style="list-style-type: none"> Read: 15,000 Write: 15,000 Black Duck Services: <ul style="list-style-type: none"> CPU: 13 core Memory: 46 GB PostgreSQL: <ul style="list-style-type: none"> CPU: 4 core Memory: 16 GB 	Total: <ul style="list-style-type: none"> CPU: 17 core Memory: 62 GB
250sph	Scans per hour: 300 SPH % Increase: 20% APIs per hour: 7,500 Project Versions: 15,000	IOPS: <ul style="list-style-type: none"> Read: 15,000 Write: 15,000 Black Duck Services: <ul style="list-style-type: none"> CPU: 17 core Memory: 118 GB PostgreSQL: <ul style="list-style-type: none"> CPU: 6 core Memory: 24 GB 	Total: <ul style="list-style-type: none"> CPU: 23 core Memory: 142 GB
500sph	Scans per hour: 650 SPH % Increase: 30% APIs per hour: 18,000 Project Versions: 18,000	IOPS: <ul style="list-style-type: none"> Read: 15,000 Write: 15,000 Black Duck Services:	Total: <ul style="list-style-type: none"> CPU: 38 core Memory: 250 GB

		<ul style="list-style-type: none"> • CPU: 28 core • Memory: 210 GB PostgreSQL: <ul style="list-style-type: none"> • CPU: 10 core • Memory: 40 GB 	
1000sph	Scans per hour: 1,400 SPH % Increase: 40% APIs per hour: 26,000 Project Versions: 25,000	IOPS: <ul style="list-style-type: none"> • Read: 25,000 • Write: 25,000 Black Duck Services: <ul style="list-style-type: none"> • CPU: 47 core • Memory: 411 GB PostgreSQL: <ul style="list-style-type: none"> • CPU: 18 core • Memory: 72 GB 	Total: <ul style="list-style-type: none"> • CPU: 65 core • Memory: 483 GB
1500sph	Scans per hour: 1,600 SPH % Increase: 6% APIs per hour: 41,000 Project Versions: 28,000	IOPS: <ul style="list-style-type: none"> • Read: 25,000 • Write: 25,000 Black Duck Services: <ul style="list-style-type: none"> • CPU: 60 core • Memory: 597 GB PostgreSQL: <ul style="list-style-type: none"> • CPU: 26 core • Memory: 104 GB 	Total: <ul style="list-style-type: none"> • CPU: 92 core • Memory: 701 GB
2000sph	Scans per hour: 2,300 SPH % Increase: 15% APIs per hour: 50,000 Project Versions: 35,000	IOPS: <ul style="list-style-type: none"> • Read: 60,000 • Write: 25,000 Black Duck Services: <ul style="list-style-type: none"> • CPU: 66 core • Memory: 597 GB PostgreSQL: <ul style="list-style-type: none"> • CPU: 34 core • Memory: 136 GB 	Total: <ul style="list-style-type: none"> • CPU: 100 core • Memory: 733 GB

This new guidance is based current Black Duck 2022.10.0 architecture. It is possible this guidance will be further refined for subsequent releases. If you have any questions or concerns, please reach out to Product Management.

 **Note:** The amount of required disk space is dependent on the number of projects being managed, so individual requirements can vary. Consider that each project requires approximately 200 MB.

Black Duck Software recommends monitoring disk utilization on Black Duck servers to prevent disks from reaching capacity which could cause issues with Black Duck.

2. Hardware requirements •

BDBA scaling is done by adjusting the number of binaryscanner replicas and by adding PostgreSQL resources based on the expected number of binary scans per hour that will be performed. For every 15 binary scans per hour, add the following:

- One binaryscanner replica
- One CPU for PostgreSQL
- 4GB memory to PostgreSQL

If your anticipated scan rate is not a multiple of 15, round up. For example, 24 binary scans per hour would require the following:

- Two binaryscanner replicas,
- Two additional CPUs for PostgreSQL, and
- 8GB additional memory for PostgreSQL.

This guidance is valid when binary scans are 20% or less of the total scan volume (by count of scans).

Binary scanning

If you are licensed for binary scanning, the uploadcache container/pod memory may need to be increased because this is where the binary scanner extracts and processes the binary. By default, the memory is set to 512MB which is not adequate for large scanning. When scanning large binaries, it is recommended to increase the memory to at least 4 GB for the uploadcache container/pod. To do so, find your override yaml and update the memory limit to 4096MB.

For Swarm installations:

```
uploadcache:
  deploy:
    resources:
      limits:
        cpus: ".200"
        memory: "4096M"
      reservations:
        cpus: ".100"
        memory: "4096M"
    replicas: 1
```

For Kubernetes installations:

```
uploadcache:
  replicas: 1
  resources:
    limits:
      cpu: "200m"
      memory: "4096Mi"
    requests:
      cpu: "100m"
      memory: "4096Mi"
```




Note: Installing Black Duck Alert requires 1 GB of additional memory.


3. PostgreSQL versions

Black Duck 2022.10.0 supports new PostgreSQL features and functionality to improve the performance and reliability of the Black Duck service. As of Black Duck 2022.10.0, PostgreSQL container 13 is the currently supported version of PostgreSQL for the internal PostgreSQL container.

Customers upgrading from older versions of Black Duck (prior to 2022.10.0), will require a migration to PostgreSQL 13. The Black Duck 2022.10.0 update migrates the internal Black Duck PostgreSQL database container to version 13 of PostgreSQL. If you use the database container and deploy on OpenShift, you need to run a one-time migration job as documented in the Black Duck release notes and installation guide.

 **Note:** For PostgreSQL sizing guidelines, see [Black Duck Hardware Scaling Guidelines](#).

If you choose to run your own external PostgreSQL instance, Synopsys recommends PostgreSQL 14 for new installs. Due to an index corruption bug in PostgreSQL 14.0 through 14.3, the minimum supported version of PostgreSQL 14 is 14.4.

 **CAUTION:** Do not run antivirus scans on the PostgreSQL data directory. Antivirus software opens lots of files, puts locks on files, etc. Those things interfere with PostgreSQL operations. Specific errors vary by product but usually involve the inability of PostgreSQL to access its data files. One example is that PostgreSQL fails with "too many open files in the system."

General Migration Process

The guidance here applies to upgrading from any PG 9.6 based Hub (releases prior to 2022.2.0) to 2022.10.0 or later.

1. The migration is performed by the blackduck-postgres-upgrader container.
2. If you are upgrading from a PostgreSQL 9.6-based Version of Black Duck:
 - The folder layout of the PostgreSQL data volume is rearranged to make future PostgreSQL version upgrades simpler.
 - The UID of the owner of the data volume is changed. The new default UID is 1001, but see the deployment-specific instructions.
3. The pg_upgrade script is run to migrate the database to PostgreSQL 13.
4. Plain ANALYZE is run on the PostgreSQL 13 database to initialize query planner statistics.
5. blackduck-postgres-upgrader exits.

4. Migrating on OpenShift with Helm

If you are upgrading from a PostgreSQL 9.6-based version of Black Duck, this migration replaces the use of a CentOS PostgreSQL container with a Synopsys-provided container. Also, the synopsys-init container is replaced with the blackduck-postgres-waiter container.

On plain Kubernetes, the container of the upgrade job will run as root unless overridden. However, the only requirement is that the job runs as the same UID as the owner of the PostgreSQL data volume (which is UID=26 by default).

On OpenShift, the upgrade job assumes that it will run with the same UID as the owner of the PostgreSQL data volume.